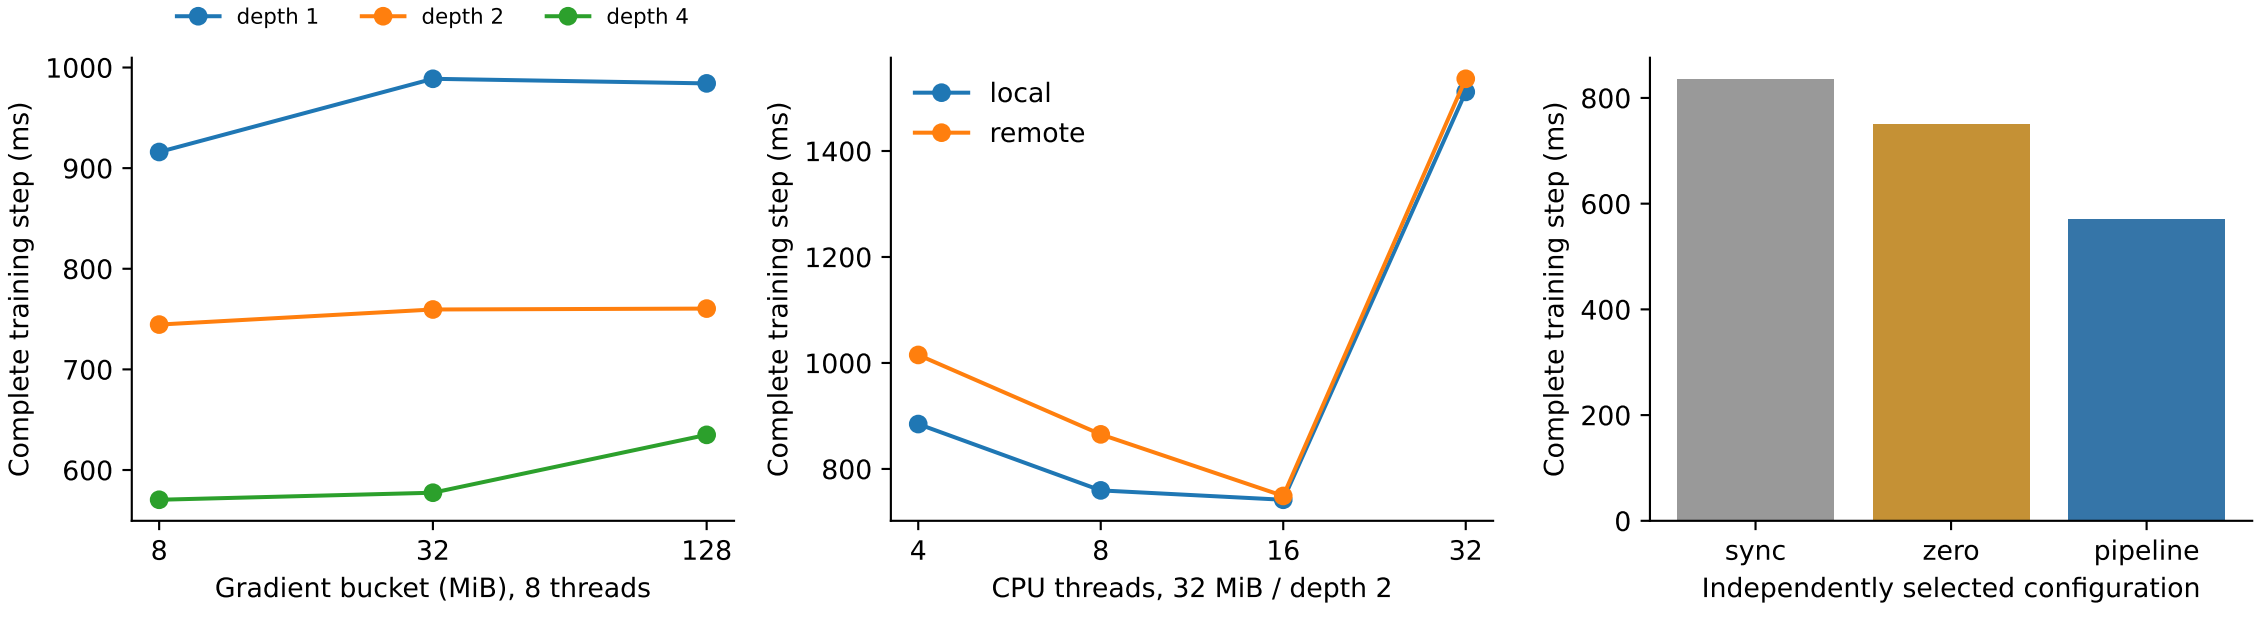
depth 1
depth 2
depth 4
Complete training step (ms)
1000
900
800
700
600
8
32
128
Gradient bucket (MiB), 8 threads
local
remote
Complete training step (ms)
1400
1200
1000
800
4
8
16
32
CPU threads, 32 MiB / depth 2
Complete training step (ms)
800
600
400
200
0
sync
zero
pipeline
Independently selected configuration